

Name : Chetan Kamatagi

College : KLE TECHNOLOGICAL UNIVERSITY, HUBBALLI, KARNATAKA

Branch : Automations and Robotics

Year : 2nd year

Github Link: <https://github.com/0311Chetan/MAJOR-PROJECT>

## MAJOR PROJECT 1:

```
# Major Project 1:
# Choose any dataset of your choice and apply a suitable CLASSIFIER/REGRESSOR.

# importing Required libraries
import pandas as pd
import numpy as np
df = pd.read_csv('https://raw.githubusercontent.com/ameenmanna8824/DATASETS/main/heart_dis
df
```

	Unnamed: 0	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	s
0	0	63	1	3	145	233	1	0	150	0	2.3	
1	1	37	1	2	130	250	0	1	187	0	3.5	
2	2	41	0	1	130	204	0	0	172	0	1.4	
3	3	56	1	1	120	236	0	1	178	0	0.8	
4	4	57	0	0	120	354	0	1	163	1	0.6	
...	...	...	...	...	...	...	...	...	...	...	...	...
298	298	57	0	0	140	241	0	1	123	1	0.2	
299	299	45	1	3	110	264	0	1	132	0	1.2	
300	300	68	1	0	144	193	1	1	141	0	3.4	
301	301	57	1	0	130	131	0	1	115	1	1.2	
302	302	57	0	1	130	236	0	0	174	0	0.0	

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 193 entries, 0 to 192
Data columns (total 25 columns):
#   Column                Non-Null Count  Dtype
---  -
0   symboling              193 non-null    int64
1   make                   193 non-null    object
2   fuel_type              193 non-null    object
3   aspiration              193 non-null    object
4   num_of_doors           193 non-null    int64
5   body_style             193 non-null    object
6   drive_wheels           193 non-null    object
7   engine_location        193 non-null    object
8   wheel_base             193 non-null    float64
9   length                 193 non-null    float64
10  width                   193 non-null    float64
11  height                 193 non-null    float64
```

```

12  curb_weight      193 non-null    int64
13  engine_type      193 non-null    object
14  num_of_cylinders  193 non-null    int64
15  engine_size      193 non-null    int64
16  fuel_system      193 non-null    object
17  bore             193 non-null    float64
18  stroke           193 non-null    float64
19  compression_ratio 193 non-null    int64
20  horsepower       193 non-null    int64
21  peak_rpm         193 non-null    int64
22  city_mpg         193 non-null    int64
23  highway_mpg      193 non-null    int64
24  price            193 non-null    int64
dtypes: float64(6), int64(11), object(8)
memory usage: 37.8+ KB

```

```

# Preprocessing of the data is not required
# because dataset is complete and there are no null values.

```

```

# Dividing the data into Input and Output

```

```

x = df.iloc[:,4:9].values

```

```

x

```

```

array([[145, 233, 1, 0, 150],
       [130, 250, 0, 1, 187],
       [130, 204, 0, 0, 172],
       ...,
       [144, 193, 1, 1, 141],
       [130, 131, 0, 1, 115],
       [130, 236, 0, 0, 174]])

```

```

y = df.iloc[:,14:15].values

```

```

y

```

```

# Train and Test variables

```

```

from sklearn.model_selection import train_test_split
xtr,xts,ytr,yts = train_test_split(x,y,random_state = 0)

```

```

print(x.shape)
print(xtr.shape)
print(xts.shape)

```

```

print(y.shape)
print(ytr.shape)
print(yts.shape)

```

```

(303, 5)
(227, 5)
(76, 5)
(303, 1)
(227, 1)
(76, 1)

```

```
# Normalizing is not necessary since the output and input are already scaled

# Running a Classifier/Regressor
from sklearn.linear_model import LogisticRegression
mdl = LogisticRegression()

# Fitting the model to the classifier:
mdl.fit(xtr,ytr)

/usr/local/lib/python3.7/dist-packages/sklearn/utils/validation.py:993: DataConversionWarning:
  y = column_or_1d(y, warn=True)
LogisticRegression()
```

```
# predicting the output:
ypred=mdl.predict(xts)
ypred

array([0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1,
       0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 0, 1, 0, 1, 1, 1, 1, 0,
       0, 0, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1,
       1, 1, 1, 1, 1, 1, 1, 0, 0, 0])
```

```
# Accuracy:
from sklearn.metrics import accuracy_score
accuracy_score(ypred,yts) *100

76.31578947368422
```

```
# Predicting the individual values
mdl.predict([[180,190,1,0,200]])

array([1])
```

```
mdl.predict([[80,90,1,0,20]])

array([0])
```

## MAJOR PROJECT 2:

```
# Major project 2:
# Choose any dataset of your choice and apply K Means Clustering.
```

```
# Import the Required Libraries
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
import numpy as np

# Importing the dataset from the github

df = pd.read_csv('https://raw.githubusercontent.com/0311Chetan/Codes/main/concrete.csv')
df
```

	Cement	BlastFurnaceSlag	FlyAsh	Water	Superplasticizer	CoarseAggregate	FineAggregate
0	540.0	0.0	0.0	162.0	2.5	1040.0	1040.0
1	540.0	0.0	0.0	162.0	2.5	1055.0	1055.0
2	332.5	142.5	0.0	228.0	0.0	932.0	932.0
3	332.5	142.5	0.0	228.0	0.0	932.0	932.0
4	198.6	132.4	0.0	192.0	0.0	978.4	978.4
...	...	...	...	...	...	...	...
1025	276.4	116.0	90.3	179.6	8.9	870.1	870.1
1026	322.2	0.0	115.6	196.0	10.4	817.9	817.9
1027	148.5	139.4	108.6	192.7	6.1	892.4	892.4
1028	159.1	186.7	0.0	175.6	11.3	989.6	989.6
1029	260.9	100.5	78.3	200.6	8.6	864.5	864.5

1030 rows × 9 columns

```
df.shape
```

```
(1030, 9)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1030 entries, 0 to 1029
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Cement                1030 non-null  float64
1   BlastFurnaceSlag      1030 non-null  float64
2   FlyAsh                1030 non-null  float64
3   Water                 1030 non-null  float64
4   Superplasticizer      1030 non-null  float64
5   CoarseAggregate       1030 non-null  float64
6   FineAggregate         1030 non-null  float64
7   Age                   1030 non-null  int64
8   CompressiveStrength   1030 non-null  float64
dtypes: float64(8), int64(1)
memory usage: 72.5 KB
```

```
# Here we take input parameters as :CoarseAggregate and FineAggregate :
```

```
# Now dividing the data into Input and output:
```

```
x = df.iloc[:,5:7].values
```


```
x
```

```
array([[1040. ,  676. ],
       [1055. ,  676. ],
       [ 932. ,  594. ],
       ...,
       [ 892.4,  780. ],
       [ 989.6,  788.9],
       [ 864.5,  761.5]])
```

```
#VISUALISATION of the input data before applying any clustering techniques :
```

```
plt.scatter(df['CoarseAggregate'],df['FineAggregate'])
```

```
<matplotlib.collections.PathCollection at 0x7fcff987b050>
```



```
#Now Finding out the number of clusters(k)
```

```
# Total number of points = 1030
```

```
np.sqrt(1030)
```

```
# k value ranges from 2 - 32
```

```
32.09361307176243
```

```
# Now we have to find No. of clusters(k)
```

```
#1.ELBOW METHOD
```

```
from sklearn.cluster import KMeans
```

```
k = range(2,33)
```

```
sse = []
```


```
for i in k :
```

```
    model_demo = KMeans(n_clusters = i,random_state = 0)
```

```
    model_demo.fit(x)
```

```
sse.append(model_demo.inertia_)
plt.scatter(k,sse)
plt.plot(k,sse)
```

[<matplotlib.lines.Line2D at 0x7fcff7f514d0>]



## # 2.SILHOUETTE SCORE METHOD

```
from sklearn.metrics import silhouette_score
k = range(2,33)
for i in k:
    model_demo = KMeans(n_clusters = i,random_state = 0)
    model_demo.fit(x)
    y_pred = model_demo.predict(x)
    print(f"{i} Clusters, Score = {silhouette_score(x,y_pred)}")
    plt.bar(i,silhouette_score(x,y_pred))
```

```

2 Clusters, Score = 0.3558185095888853
3 Clusters, Score = 0.409974801365556
4 Clusters, Score = 0.4130838519874459
5 Clusters, Score = 0.4486948545137604
6 Clusters, Score = 0.45199448215467036
7 Clusters, Score = 0.470588180756573
8 Clusters, Score = 0.46567284221181254
9 Clusters, Score = 0.4636699148596504
10 Clusters, Score = 0.4691868710743738
11 Clusters, Score = 0.46267961260749596
12 Clusters, Score = 0.4658701556028496
13 Clusters, Score = 0.4885732054608967
14 Clusters, Score = 0.4828828733854631
15 Clusters, Score = 0.46943457773884695
16 Clusters, Score = 0.4796691111373154
17 Clusters, Score = 0.4840973594478049
18 Clusters, Score = 0.47860901217317503
19 Clusters, Score = 0.48117825012833304
20 Clusters, Score = 0.48395282373636606
21 Clusters, Score = 0.4972722120915127
22 Clusters, Score = 0.4847935953371792
23 Clusters, Score = 0.5018150406160546

```

# From the SILHOUETTE SCORE METHOD we can clearly observe that there can be 32 clusters ca  
 # THE No OF CLUSTERS TO BE CONSIDERED IS 32.

```

28 Clusters, Score = 0.5280081121806302

```

# 7.APPLY CLUSTERER

k = 32

from sklearn.cluster import KMeans

model = KMeans(n\_clusters = k, random\_state = 0)

model.fit(x)

```

KMeans(n_clusters=32, random_state=0)

```

```

|  0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 |

```

y = model.predict(x) # predicted output

y

```

array([31, 31, 10, ..., 1, 2, 1], dtype=int32)

```

```

|  0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 |

```

y.size

```

1030

```

x[y == 1,1]

#so the first '1' is cluster no 1 and the second '1' is column index 1

#the value of input,when cluster 1 is selected and column index 1 selected

```

array([781.5, 781.5, 781.5, 781.5, 781.5, 805. , 768. , 800. , 750. ,
       785. , 774. , 790. , 780. , 768. , 780. , 761. , 783. , 785.3,
       774. , 790. , 779.7, 768.3, 780. , 761.5])

```

np.unique(y,return\_counts = True)



```
(array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
        17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31],
      dtype=int32),
 array([68, 24, 35,  7, 51, 39, 44, 12, 31, 31, 32, 56, 20, 37, 40, 20, 31,
        87, 52, 21, 15, 22, 64, 23, 43,  9, 10, 21,  5, 37, 28, 15]))
```

```
#FINAL VISUALISATION
```

```
plt.figure(figsize = (10,5))
```


```
for i in range(k):
```

```
    plt.scatter(x[y == i,0],x[y == i,1],label = f'Cluster {i}')
```

```
plt.scatter(model.cluster_centers_[0],model.cluster_centers_[1],s = 300,c = 'yellow',
            label = 'Centroids')
```

```
plt.legend()
```

```
<matplotlib.legend.Legend at 0x7fcff7b975d0>
```



[Colab paid products](#) - [Cancel contracts here](#)

---

✓ 1s    completed at 10:37 PM ● ✕



Reconnect  Editing 




```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 193 entries, 0 to 192
Data columns (total 25 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   symboling                             193 non-null    int64
1   make                                  193 non-null    object
2   fuel_type                             193 non-null    object
3   aspiration                             193 non-null    object
4   num_of_doors                          193 non-null    int64
5   body_style                            193 non-null    object
6   drive_wheels                          193 non-null    object
7   engine_location                       193 non-null    object
8   wheel_base                            193 non-null    float64
9   length                               193 non-null    float64
10  width                                 193 non-null    float64
11  height                               193 non-null    float64
12  curb_weight                           193 non-null    int64
13  engine_type                           193 non-null    object
14  num_of_cylinders                      193 non-null    int64
15  engine_size                           193 non-null    int64
16  fuel_system                           193 non-null    object
17  bore                                  193 non-null    float64
18  stroke                                193 non-null    float64
19  compression_ratio                     193 non-null    int64
20  horsepower                             193 non-null    int64
21  peak_rpm                              193 non-null    int64
22  city_mpg                              193 non-null    int64
23  highway_mpg                           193 non-null    int64
24  price                                 193 non-null    int64

dtypes: float64(6), int64(11), object(8)
memory usage: 37.8+ KB
```

✓ 1s completed at 10:37 PM



 Data Science Projects.ipynb ☆

File Edit View Insert Runtime Tools Help Save failed

+ Code + Text

Reconnect Editing

```
# Running a Classifier/Regressor
from sklearn.linear_model import LogisticRegression
mdl = LogisticRegression()

# Fitting the model to the classifier:
mdl.fit(xtr,ytr)

/usr/local/lib/python3.7/dist-packages/sklearn/utils/validation.py:993: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to
y = column_or_1d(y, warn=True)
LogisticRegression()

[ ] # predicting the output:
ypred=mdl.predict(xts)
ypred

array([[0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1,
        0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 0, 1, 0, 1, 1, 1, 1, 0,
        0, 0, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 0,
        0, 0, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 1, 1, 0, 1, 0, 1,
        1, 1, 1, 1, 1, 1, 1, 0, 0, 0]])

[ ] # Accuracy:
from sklearn.metrics import accuracy_score
accuracy_score(ypred,yts) *100

76.31578947368422

[ ] # Predicting the individual values
mdl.predict([[180,190,1,0,200]])

array([1])

[ ] mdl.predict([[80,90,1,0,20]])

array([0])
```

✓ 1s completed at 10:37 PM



Reconnect  Editing 

(1030, 9)

```
>>> <class 'pandas.core.frame.DataFrame'>
RangeIndex: 1030 entries, 0 to 1029
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   cement                                1030 non-null   float64
1   BlastFurnaceSlag                     1030 non-null   float64
2   FlyAsh                                1030 non-null   float64
3   Water                                 1030 non-null   float64
4   Superplasticizer                     1030 non-null   float64
5   CoarseAggregate                      1030 non-null   float64
6   FineAggregate                        1030 non-null   float64
7   Age                                   1030 non-null   int64
8   CompressiveStrength                  1030 non-null   float64
dtypes: float64(8), int64(1)
memory usage: 72.5 KB
```

```
array([[1040. , 676. ],
       [1055. , 676. ],
       [ 932. , 594. ],
       ...,
       [ 892.4, 780. ],
       [ 989.6, 788.9],
       [ 864.5, 761.5]])
```




Reconnect  Editing 

```
plt.scatter(df['CoarseAggregate'],df['FineAggregate'])
```


```
sse = []
```

Reconnect   Editing 

```
[<matplotlib.lines.Line2D at 0x7fcff7f514d0>]
```



✓ 1s completed at 10:37 PM

 Data Science Projects.ipynb ☆

File Edit View Insert Runtime Tools Help Save failed

+ Code + Text

Reconnect Editing


```
# 2.SILHOUETTE SCORE METHOD

from sklearn.metrics import silhouette_score
k = range(2,33)
for i in k:
    model_demo = KMeans(n_clusters = i,random_state = 0)
    model_demo.fit(x)
    y_pred = model_demo.predict(x)
    print(f"{i} Clusters, Score = {silhouette_score(x,y_pred)}")
    plt.bar(i,silhouette_score(x,y_pred))
```

```
2 Clusters, Score = 0.3558185095888853
3 Clusters, Score = 0.409974801365556
4 Clusters, Score = 0.4130838519874459
5 Clusters, Score = 0.4486948545137604
6 Clusters, Score = 0.45199448215467036
7 Clusters, Score = 0.470588180756573
8 Clusters, Score = 0.46567284221181254
9 Clusters, Score = 0.4636699148596504
10 Clusters, Score = 0.4691868710743738
11 Clusters, Score = 0.46267961260749596
12 Clusters, Score = 0.4658701556028496
13 Clusters, Score = 0.4885732054608967
14 Clusters, Score = 0.4828828733854631
15 Clusters, Score = 0.4694345773884695
16 Clusters, Score = 0.4796691111373154
17 Clusters, Score = 0.4840973594478049
18 Clusters, Score = 0.4786090121717503
19 Clusters, Score = 0.4811782501283304
20 Clusters, Score = 0.48395282373636606
21 Clusters, Score = 0.4972722120915127
22 Clusters, Score = 0.4847935953371792
23 Clusters, Score = 0.5018159496160546
24 Clusters, Score = 0.5098744337479676
25 Clusters, Score = 0.5105166834519095
26 Clusters, Score = 0.5136563831086176
27 Clusters, Score = 0.5133755643171831
28 Clusters, Score = 0.5280081121806302
```

Automatic document saving has been pending for 6 minutes. Reloading may fix the problem. [Save and reload the page.](#)

10 completed at 10:07 PM



```
[ ] # 7. APPLY CLUSTERER
k = 32
from sklearn.cluster import KMeans

model = KMeans(n_clusters = k, random_state = 0)
model.fit(x)

KMeans(n_clusters=32, random_state=0)
```



